

PULP: Parallel Ultra Low Power a platform for next generation IoT Applications

ACA PROJECT 2016/2017
Prof. Cristina Silvano

Stefano Brandoli



POLITECNICO
MILANO 1863

Introduction: IoT world

REALITY:

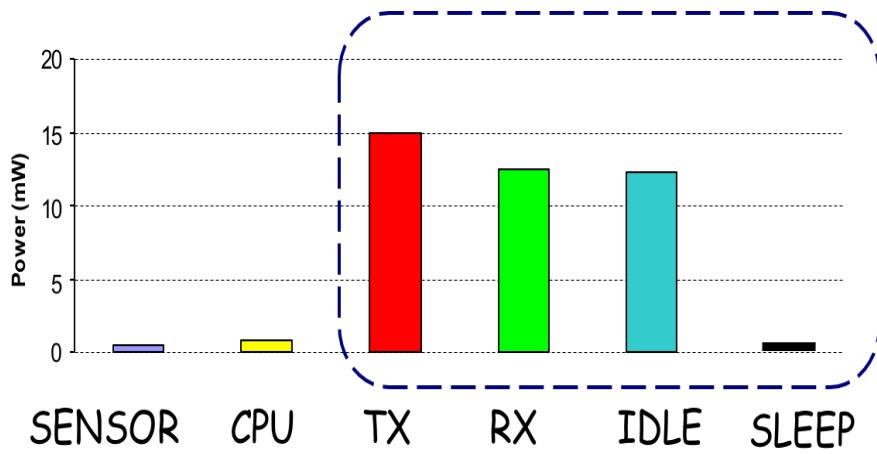
Sensor nodes have limited power source (batteries)

PROBLEM:

Recharging and/or battery replacement may be immaterial or too expensive

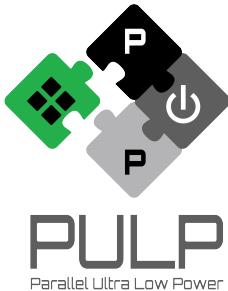
SOLUTIONS:

- Networking Protocols
- Low-power processing
- Low-power processors
- Low-power sensing
- Energy Harvesting
- Tx/Rx consume much of the power
 - disable them as soon as possible
- Idle consumes similarly to Tx/Rx
 - to reduce power consumption we have to reduce IDLE time
- Low CPU Power consumption
 - sensor node does simple tasks



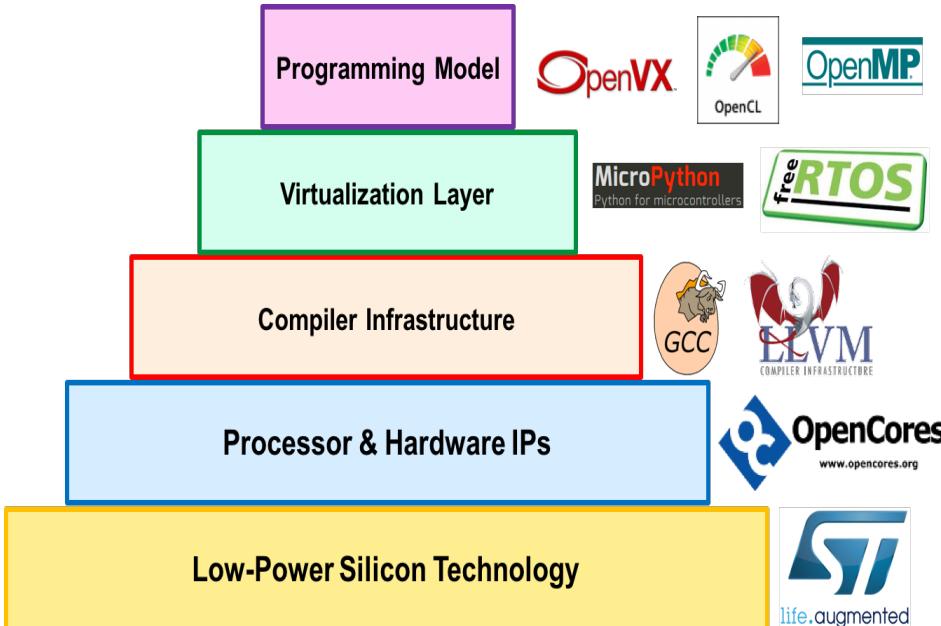
GOAL: ENABLE MORE COMPLEX IOT APPLICATIONS USING THE SAME POWER BUDGET (few mW)

- Performance: architecture parallelism
- Energy Efficiency: voltage scaling without losing performance (**near threshold operations**)



Introduction: The PULP platform

- Energy efficient **many core SoC**
- **Software frameworks** (*OpenMP, OpenCL, OpenVX*):
 - exploit hardware parallelism
 - enable agile portings, development, debugging, tuning
- **Custom toolset** (*Virtual Platform, profiling tools, FPGA emulation...*)
- Large number of IPs



pJ/op = target of uControllers and ASIC devices

- not able to deliver the required computational power

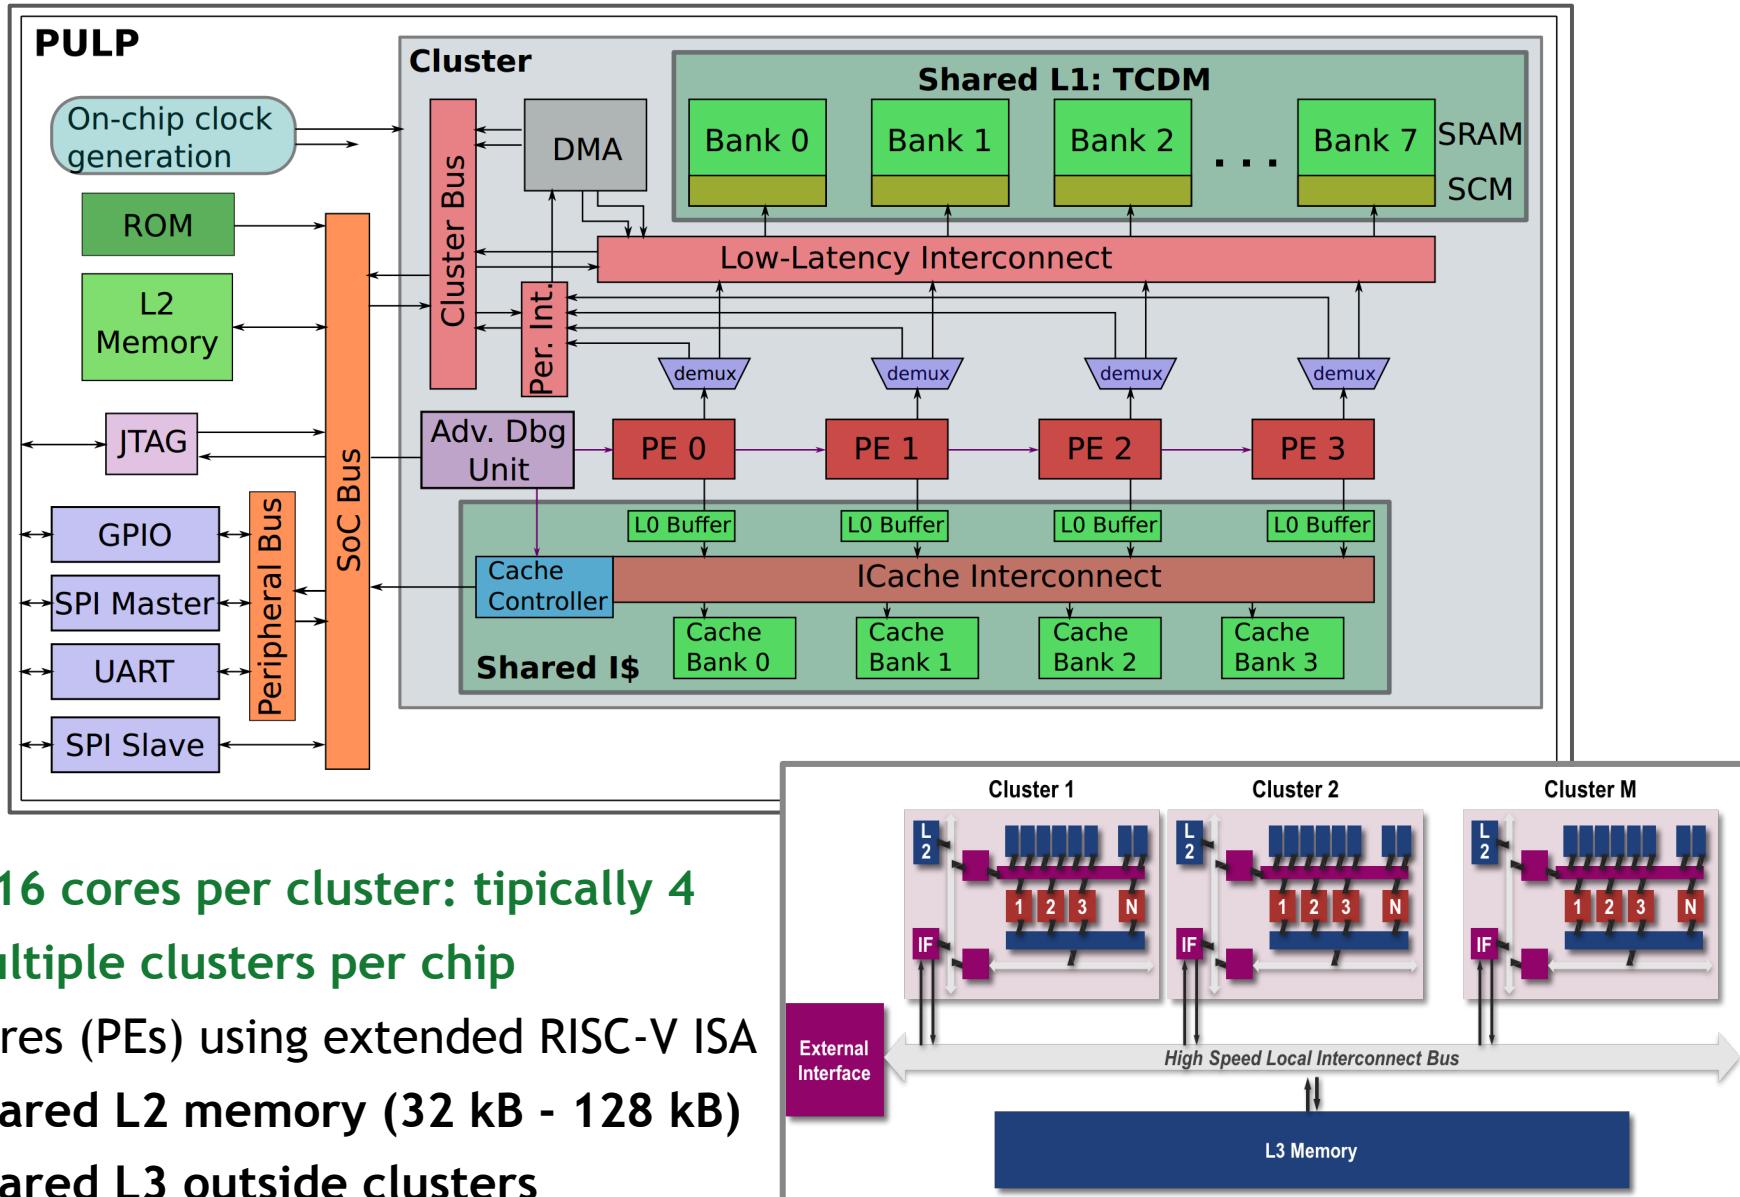
IDEA:

- expose low power features of the technology
- use low power silicon technology

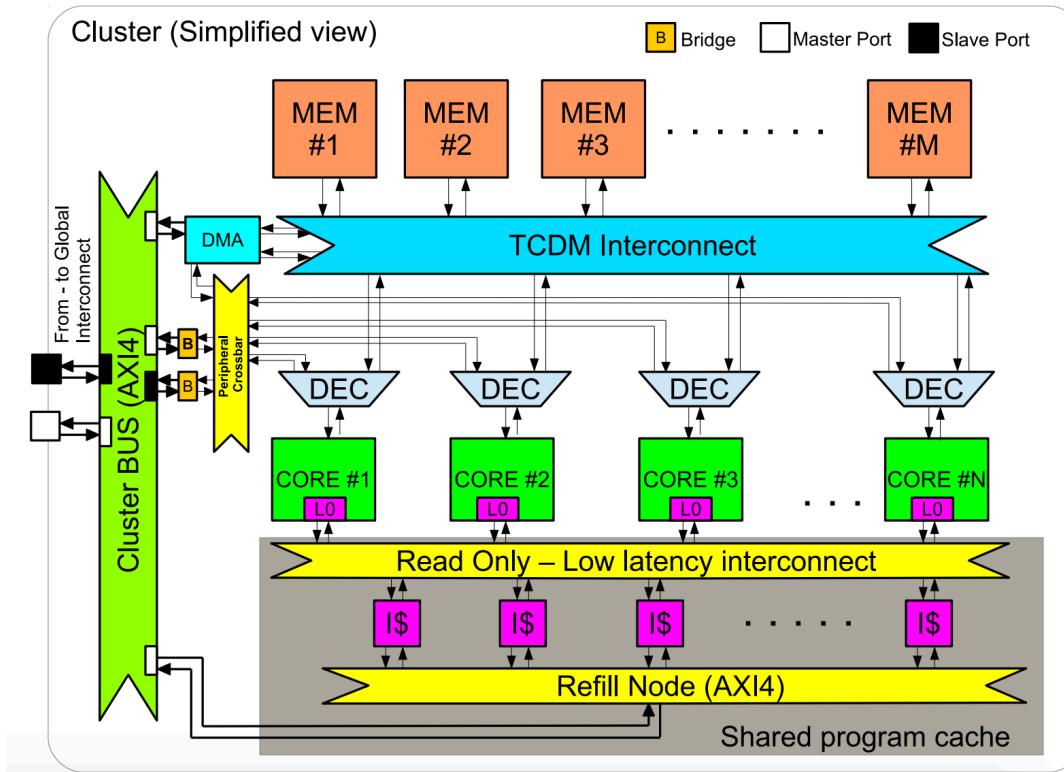
FEATURES:

- **67 MOPS/mW**: 65nm bulk CMOS
- **193 MOPS/mW** (40MHz, 1mW): low power STM 28nm FDSOI

Architecture: Overview



Architecture: Shared I-Cache



- The **Refill Node (AXI4)** of the shared I-Cache connects to the system bus
- System bus connects to:
 - L2 memory
 - Several peripherals
- Shared I-Cache implemented with SCMs

Architecture: I-Cache

- Shared 4-way set associative I-Cache with 4 cache banks of 1 KB each
- Shared I-Cache vs Private I-Cache

- **Avoid Data Replication:**

several cores fetch the same code segment, that is stored in each **private I-cache** (copies). Multiple refill requests are asserted at the same address, creating intense traffic to a slow and far L2 memory subsystem (~10-40 cycles)

- **Better Miss Ratio:**

bigger memory space seen by each core

- **Less speed**

more complex design (vs Direct Mapped)

- **Increased HW complexity**

- **L0 Pre-fetch Buffer (128 bit per core):**

- Contains a complete cache line (128 bit)

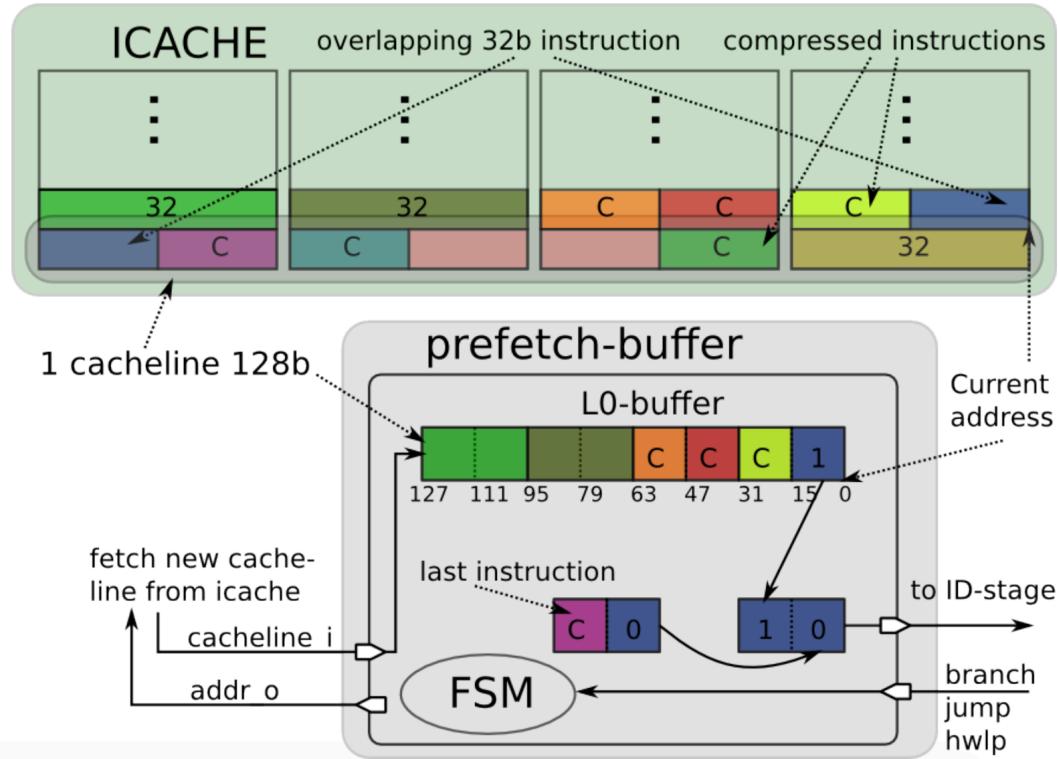
- **Reduce pressure (and contention) on shared I-Cache:**

- only in case of miss the request is forwarded to the shared I-Cache (in the same cycle)

PROs

CONs

Architecture: L0 Pre-Fetch Buffer



Support for Compressed Instruction (16 bit):

Inevitably some 32 bit instructions will become unaligned.

Requires an additional fetch cycle for which the processor has to be stalled

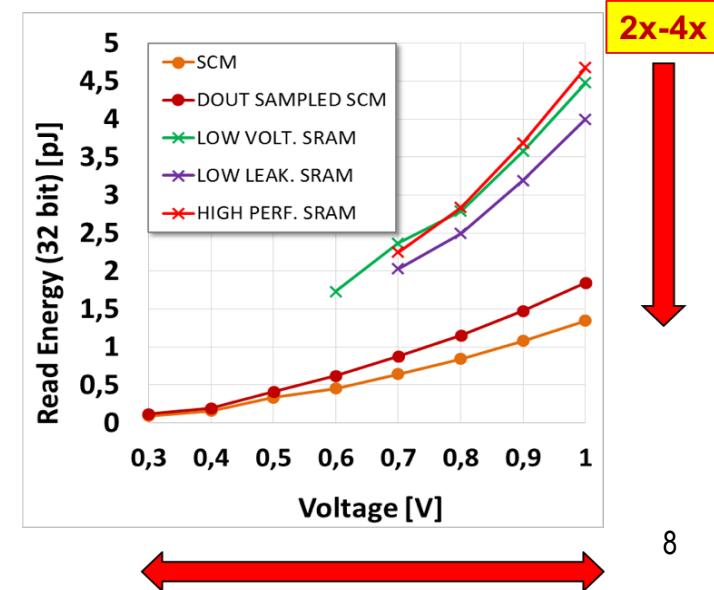
- **SOLUTION: use an additional register to keep the last instruction:**

this register will contain the lower 16 bit of the instruction which can be combined with the higher part and forwarded to the ID-stage

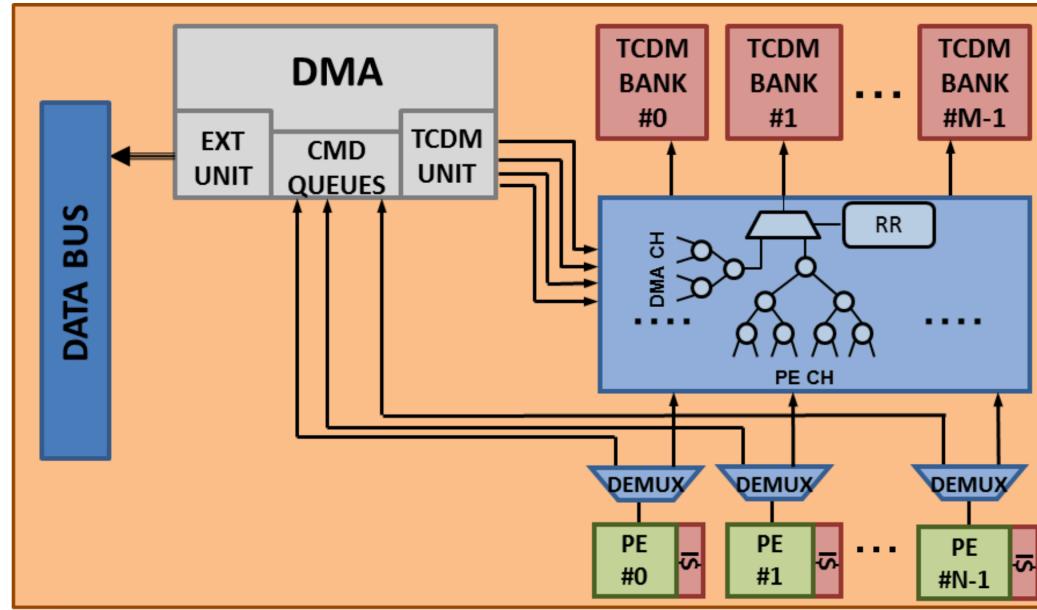
Architecture: Shared Data Memory

- Cores don't have private Data Caches:
 - avoid memory coherence overhead
 - improving power and area efficiency for data memory
- Shared L1 multi-banked TCDM (*Tightly Coupled Data Memory*):
 - parametric number of word-level interleaved banks
 - acts as a software-managed shared data scratchpad
 - logical banks connected to the processors through a non-blocking interconnect to minimize banking conflicts
 - 1 cycle communication between PEs and memory banks (read/write)
 - 1 port for every memory bank
 - read unaligned data in only 2 cycles
- TCDM logical banks are heterogeneous memory:
 - SRAMs (16x4 KB banks):
 - higher density than SCMs (3x-4x)
 - **performance degrades at low voltage**
 - low technology portability
 - Standard Cell Memory (SCM) banks (16x0.5 KB):
 - work over the same voltage range as logic
 - energy/access lower than that of SRAMs
 - high technology portability

256x32 6T SRAMS vs. SCM



General Architecture: DMA



GOAL: overlap as much as possible computation and memory transfers

DMA:

- Enables communication with other clusters, L2 memory and external peripherals
- Hides latencies of slow peripherals and of L2 memory data transfers
- **DMA + TCDM -> Double Buffering**

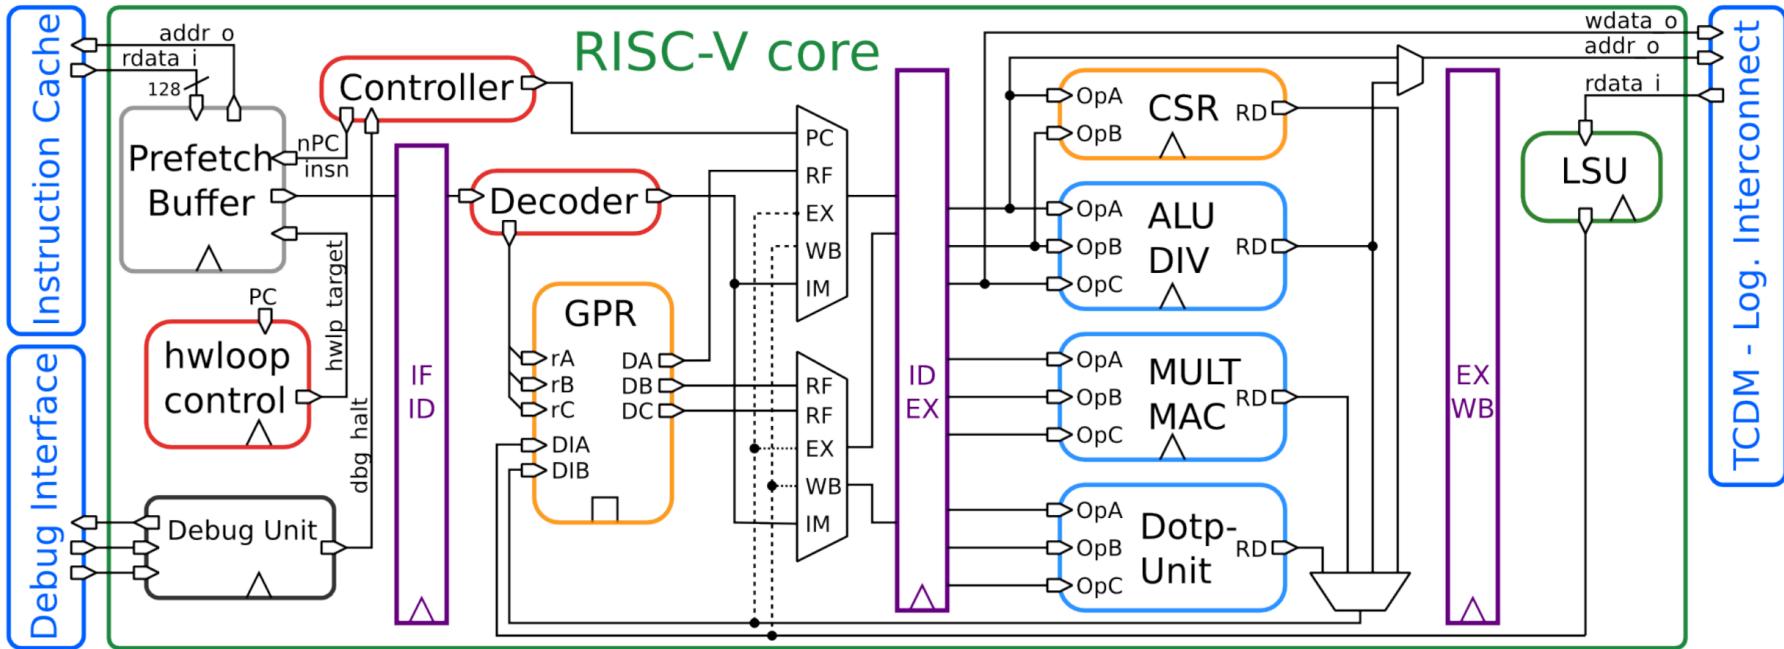
General Architecture: DMA

- **Lock-free per-core FIFO command queues:**
 - Each processor has the ability to configure DMA transfers by sending commands to a dedicated memory mapped control interface
 - If the command queue is full, the incoming request is stalled and it is committed when one location is released
 - Enqueuing and arbitration of transfer requests from PEs is performed internally by the DMA. This process is transparent to the software
- **No large internal data buffers:**
 - exploiting the cluster's shared memory as local repository
- **External Unit (EXT UNIT):**
 - decoupled read and write channels to optimize the performance when concurrent TX and RX transactions are required
 - parametric number of outstanding transactions (up to 16)
 - handles on-the-fly reordering of incoming packets
- **TCDM Unit:**
 - 4 (2 read, 2 write) ports (32-bit) for the connection with the TCDM interconnect
 - **TCDM interconnect is implemented as two separate trees**, one handling requests from processors and one used for the DMA: **Round Robin arbitration**

Power Management

- Each PULP cluster and the rest of the SoC are in different power and clock domains (**DVFS - Dynamic voltage and frequency scaling**)
- Each processor within a cluster can be separately **disabled, clock-gated and reverse body biased when idle**
- The cluster is able to keep active an arbitrary number of processing elements, while the others consume zero dynamic power and up to 10x less leakage power
- An **event unit (HW SYNCRONIZER)** automatically manages transitions of the cores between the active and idle states
- Processors can be put in idle state with a write operation on a memory mapped control register. After going in sleep mode, a core remains idle until a configurable event is triggered
- Events can be issued by all IO peripherals, the DMA, and timers

RI5CY: Optimized RISC-V core



- Shallow 4-stage in-order 32b RISC-V processor core
- Comparable to an ARM Cortex M4 in terms of area and performance
- Focus on energy efficiency: allow the core to sleep and wait for an event

RI5CY: RISC-V ISA + Custom Extensions

- Full support for RV32I Base Integer Instruction Set
 - Full support for RV32C Standard Extension for Compressed Instructions
 - Smaller code size, less pressure on instruction cache
 - Requires support for 16-bit aligned memory access
 - Full support for RV32M Integer Multiplication and Division Instruction Set
- PULP specific extensions:
- Hardware Loops
 - Post-Incrementing load and stores
 - ALU extensions (min, max, abs, avg)
common in IoT scenarios
 - Multiply-Accumulate ($d = a \times b + c$)
MAC with three input operands
 - PACKET SIMD:
 - example: Dot Product: instructions that can multiply four pairs of 8 bit numbers and accumulate them in a single instruction

RI5CY Extension: Hardware Loops

- Execute a piece of code multiple times, without the overhead of branches or updating a counter
- 0 stall cycles for jumping to the first instruction of a loop
- Start address, end address, counter -> registers mapped in CSR address space
- Max 2 nested loops
- Minimum loop size is two instructions
- Loop counter > 0: the loop body is always entered at least once

Simple vector addition

```
for (i = 0; i < 100; i++) {  
    d[i] = a[i] + 1;  
}
```

Baseline

```
        mv    x4, 0  
        mv    x5, 100  
Lstart: lw    x2, 0(x10)  
         addi x10, x10, 4  
         addi x2, x2, 1  
         sw    x2, 0(x11)  
         addi x11, x11, 4  
         addi x4, x4, 1  
         bne   x4, x5, Lstart
```

with hardware loops

```
lp.setupi 100, Lend  
lw    x2, 0(x10)  
addi x10, x10, 4  
addi x2, x2, 1  
sw    x2, 0(x11)  
Lend: addi x11, x11, 4
```

7 instr. + branch cost

5 instr.

RI5CY Extension: Post-incrementing LD/ST

- Perform a load/store operation from/to the data memory while at the same time increasing the base address by the specified offset
- Reduce the number of required instructions to execute code with regular data access patterns (like loops)
- ALU and LSU have a separate RF WRITE PORT to support this ISA extension

Simple vector addition

```
for (i = 0; i < 100; i++) {  
    d[i] = a[i] + 1;  
}
```

with hardware loops

```
lp.setupi 100, Lend  
lw    x2, 0(x10)  
addi x10, x10, 4  
addi x2, x2, 1  
sw    x2, 0(x11)  
Lend: addi x11, x11, 4
```

5 instr.

with hardware loops & post-incr.

```
lp.setupi 100, Lend  
lw    x2, 4(x10!)  
addi x2, x3, 1  
Lend: sw    x2, 4(x11!)
```

3 instr.

References: Papers

- Davide Rossi, Francesco Conti, Andrea Marongiu, Antonio Pullini, Igor Loi, Michael Gautschi, Giuseppe Tagliavini, Philippe Flatresse, Luca Benini
PULP: A Parallel Ultra-Low-Power Platform for Next Generation IoT Applications
<https://drive.google.com/file/d/0B08HOJaoGEHSQXd2UmpkU1lSVTQ>
- Michael Gautschi, Pasquale Davide Schiavone, Andreas Traber, Igor Loi, Antonio Pullini, Davide Rossi, Eric Flamand, Frank K. Gürkaynak, Luca Benini
A near-threshold RISC-V core with DSP extensions for scalable IoT Endpoint Devices
<https://arxiv.org/pdf/1608.08376.pdf>
- Igor Loi, Davide Rossi, Germain Haugou, Michael Gautschi, Luca Benini
Exploring Multi-banked Shared-L1 Program Cache on Ultra-Low Power Tightly Coupled Processor Clusters
<http://dl.acm.org/citation.cfm?id=2747288>
- Davide Rossi, Igor Loi, Germain Haugou, Luca Benini
Ultra-low-latency lightweight DMA for tightly coupled multi-core clusters
<https://drive.google.com/open?id=0B08HOJaoGEHSwUtINWxYYXkzZEK>
- Michael Gautschi, Andreas Traber, Antonio Pullini, Luca Benini, Michele Scandale, Alessandro Di Federico, Michele Beretta, Giovanni Agosta
Tailoring instruction-set extensions for an ultra-low power tightly-coupled cluster of OpenRISC cores
<https://drive.google.com/open?id=0B08HOJaoGEHSYV9SYOFDSnNjVnc>

References: General

- <http://iis-projects.ee.ethz.ch/index.php/PULP>
- <http://www.pulp-platform.org/>
- <http://home.deib.polimi.it/cesana/teaching/IoT/2017/classes/2.Hardware.pdf>
- https://riscv.org/wp-content/uploads/2016/01/Wed1315-PULP-riscv3_noanim.pdf
- http://www.pulp-platform.org/wp-content/uploads/2017/02/ri5cy_user_manual.pdf
- https://www.dropbox.com/s/qs3jbqqqiz0948tj/PULP_ORCONF15.pptx?dl=0
- http://www.hotchips.org/wp-content/uploads/hc_archives/hc27/HC27.24-Monday-Epub/HC27.24.10-IoT-Epub/HC27.24.111-PULP-Rossi-DEL-ETH-2.pdf